

TQS Lab activities

v2020-04-30

Lab 1: Unit testing with JUnit 5	2
Learning objectives	2
Preparatory readings	2
Key points	2
Lab activities	2
Explore	4
Lab 2: Mocking dependencies in unit testing	5
Learning objectives	5
Lab activities	5
Explore	7
Lab 3: Functional testing with web automation	7
Prepare	7
Key Points	7
Lab	7
Explore	8
Lab 4: Behavior-driven development (Cucumber in Java)	9
Key Points	9
Lab	9
Lab 5: Multi-layer application testing with Spring Boot	10
Prepare	10
Key Points	10
Lab	10
Lab 6: Static Code analysis with Sonar Qube	13
Prepare	13
Key Points	13
Lab	14
Explore:	16
Lab 7: Continuous integration with Jenkins	16
Prepare	16
Lab	16
Lab 8: Code review in the context of pull/merge requests	20
Prepare	20
Lab	20
Explore	21
Lab 9: Performance testing	21
Key points	21
Lab	21
Explore	22

Lab 1: Unit testing with JUnit 5

Learning objectives

- Identify relevant unit tests to verify the contract of a module.
- Write and execute unit tests using the JUnit framework.
- Link the unit tests results with further analysis tools (e.g.: code coverage)

Preparatory readings

- The [test pyramid concept](#).
- Optional: [TDD & Unit testing in IntelliJ](#) tutorial.

Key points

- Unit testing is when you (as a programmer) write test code to verify units of (production) code. A unit represents some small subset of a much larger end-to-end-behavior. A true “unit” does not have dependencies on other (external) components.
- Unit tests help the developers to (i) understand the module contract (what to construct); (ii) document the intended use of a component; (iii) prevent regression errors; (iv) increase confidence on the code.
- When following a TDD approach, typically you go through a cycle of [Red-Green-Refactor](#). You’ll run a test, see it fail (go red), implement the simplest code to make the test pass (go green), and then refactor the code so your test stays green and your code is sufficiently clean.
- JUnit and TestNG are popular frameworks for unit testing in Java.

JUnit best practices: unit test one object at a time

A vital aspect of unit tests is that they’re finely grained. A unit test independently examines each object you create, so that you can isolate problems as soon as they occur. If you put more than one object under test, you can’t predict how the objects will interact when changes occur to one or the other. When an object interacts with other complex objects, you can surround the object under test with predictable test objects. Another form of software test, integration testing, examines how working objects interact with each other. See chapter 4 for more about other types of tests.

Lab activities

Be sure that your developer environment meets the following requirements:

- Java development environment ([JDK](#)), v8 or v11. Note that you should install it into a path without spaces or special characters (e.g.: avoid \Users\José Conceição\Java).
- [Maven configured](#) to run in the command line.
- Java capable IDE, such as [IntelliJ IDEA](#).

Implement a stack data structure (TqsStack) with appropriate unit tests. Be sure to adopt a **write-the-tests-first** workflow:

- a) Create a new **maven-based**, Java standard application.

Note: use the IDE features. If you are not sure if the IDE can generate a maven-compatible structure, consider using this [starter project](#).

- b) Create the required classes definition (**just the “skeleton”**, not the methods body; you may need to add dummy return values). The code should compile, though the implementation is incomplete.
- c) Write the unit tests that will verify the TqsStack contract.

You may use the IDE features to generate the testing class; note that the [IDE support will vary](#). Be sure to use [JUnit 5.x](#).

Important: for [maven based projects](#), check the proper POM.xml [dependencies for JUnit 5](#). If you get an error saying that the 1.5 version is no longer supported, specify the [target compiler version](#) in the POM.

Your tests will verify several [assertions that should evaluate to true](#) for the test to pass. See [some examples](#).

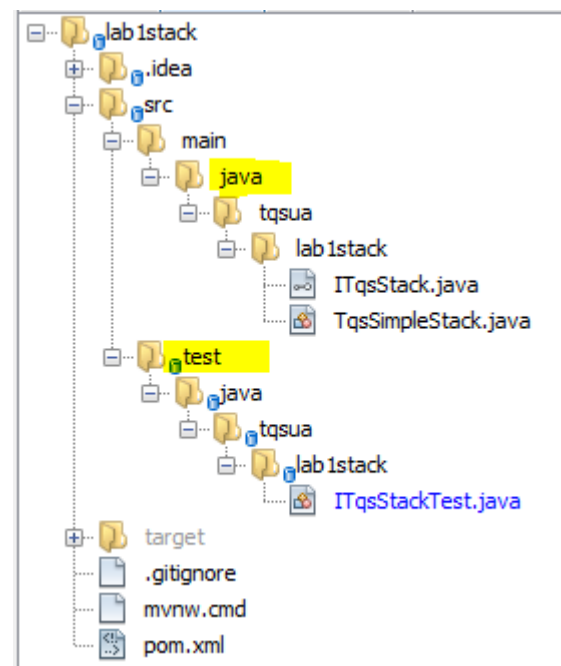
- d) Run the tests and prove that TqsStack implementation is not valid yet (the tests should fail for now, the first step in [Red-Green-Refactor](#)).
- e) Correct/add the missing implementation to the TqsStack;
- f) Run the unit tests.
- g) Iterate from steps d) to f) and confirm that all tests pass.

Stack operations:

- push(x): add an item on the top
- pop: remove the item at the top
- peek: return the item at the top (without removing it)
- size: return the number of items in the stack
- isEmpty: return whether the stack has no items

What to test¹:

- a) A stack is empty on construction.
- b) A stack has size 0 on construction
- c) After n pushes to an empty stack, $n > 0$, the stack is not empty and its size is n
- d) If one pushes x then pops, the value popped is x.
- e) If one pushes x then peeks, the value returned is x, but the size stays the same
- f) If the size is n, then after n pops, the stack is empty and has a size 0
- g) Popping from an empty stack does throw a NoSuchElementException [[You should test for the Exception occurrence](#)]
- h) Peeking into an empty stack does throw a NoSuchElementException
- i) For bounded stacks only, pushing onto a full stack does throw an IllegalStateException



2a/ Pull the [“euromillions-play” project](#) and correct the code (or the tests themselves, if needed) to have the existing unit tests passing.

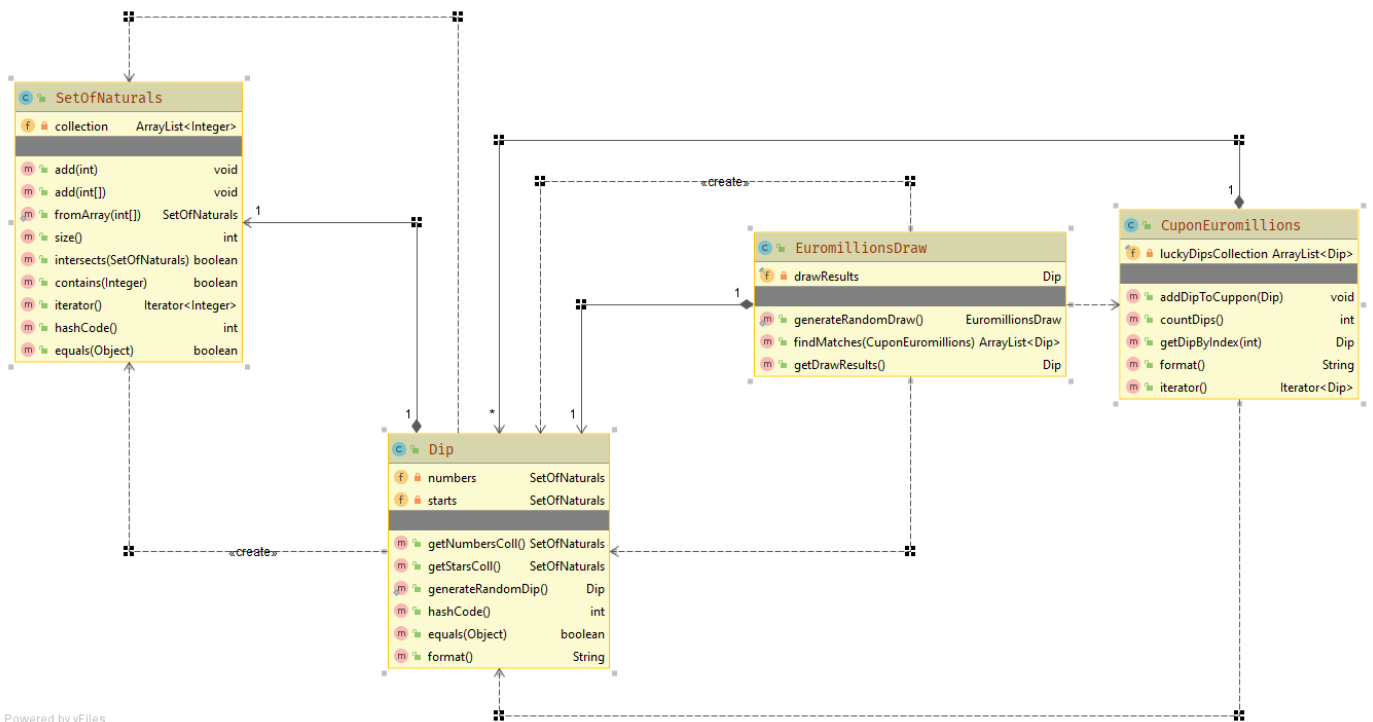
For test:	You should:
testFormat	Correct the <u>implementation</u> of Dip#format so the tests pass.

¹ Adapted from <http://cs.lmu.edu/~ray/notes/stacks/>

testConstructorFromBadArrays

Implement new test logic to confirm that an exception will be raised if the arrays have invalid numbers (wrong count of numbers of starts)

Note: you may suspend temporarily a test with the `@Disabled` tag (useful while debugging the tests themselves).



Powered by yFiles

2b/ The class `SetOfNaturals` represents a set (no duplicates should be allowed) of integers, in the range $[1, +\infty]$. Some basic operations are available (add element, find the intersection...).

What kind of unit test are worth writing for the entity `SetOfNaturals`?

Complete the project, adding the new tests you identified.

2c/ Note that the provided code includes “magic numbers” (2 for the number of stars, 50 for the max range,...). Refactor the code to extract constants and eliminate the “magic numbers” bad-smell.

2d/ Assess the coverage level in project “Euromillions-play”.

[Configure the maven project to run Jacoco analysis.](#)

Run the maven “test” goal and then “jacoco:report” goal. You should get an HTML report under `target/jacoco`.

Interpret the results accordingly. Which classes/methods offer less coverage? Are all possible decision branches being covered?

Note: IntelliJ has an integrated option to run the tests with the coverage checks (without setting the Jacoco plugin in POM). But if you do it at maven level, you can use this feature in multiple tools.

Explore

- Book: [JUnit in Action](#).
- Vogel's [tutorial on JUnit](#). Useful to compare between JUnit 4 and JUnit 5.
- [Working effectively with unit testing](#) (podcast).

Lab 2: Mocking dependencies in unit testing

Learning objectives

- Prepare a project to run unit tests ([JUnit 5](#)) and mocks ([Mockito 3.x](#)), with mocks injection (@Mock).
- Write and execute unit tests with mocked dependencies.
- Play with mock behaviors: strict/lenient verifications, advanced verifications, etc.

Lab activities

1a/ Implement the test case illustrated with the following classes, with respect to the **StockPortfolio#getTotalValue()** method. The method is expected to calculate the value of the portfolio by summing the current value (looked up in the stock market) of the owned stocks. Be sure to use:

- Maven-based Java application project;
- Mockito framework (mind the maven dependencies).

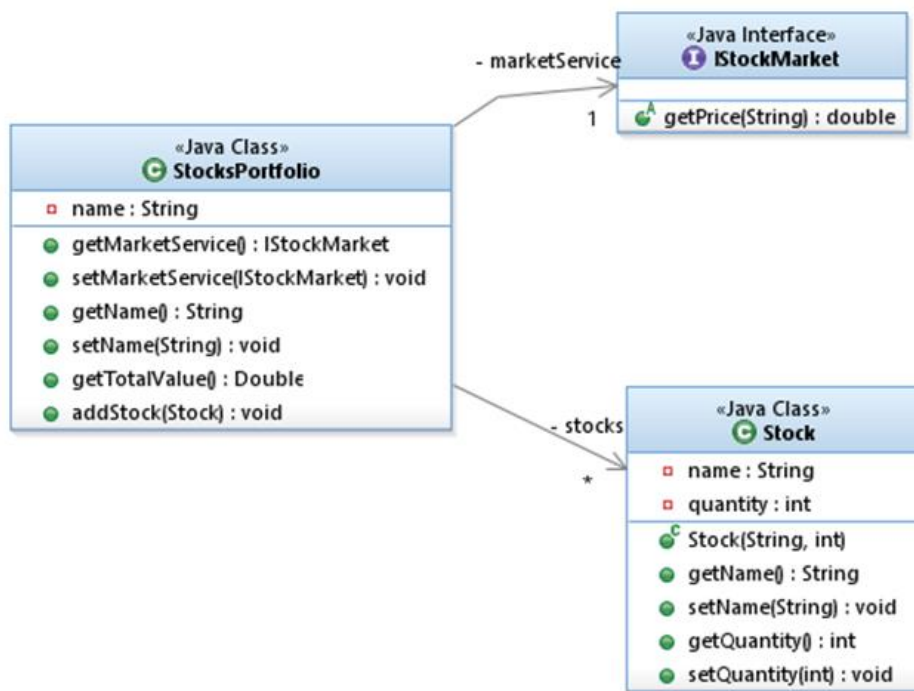


Figure 1: Classes for the StocksPortfolio use case.

- Create the classes. You may write the implementation of the services before or after the tests.
- Create the test for the `getTotalValue()`. As a guideline, you may adopt this outline:
 - Prepare a mock to substitute the remote service (@Mock annotation)
 - Create an instance of the subject under test (SuT) and use the mock to set the (remote) service instance (you may prefer to use @InjectMocks)
 - Load the mock with the proper expectations (when...thenReturn)
 - Execute the test (use the service in the SuT)

5. Verify the result (assert) and the use of the mock (verify)

Notes:

- Mind the JUnit version. For JUnit 5, you should use the `@ExtendWith` annotation to integrate the Mockito framework.
- Some IDE may not support JUnit 5 integration; you may need to [further configure the POM](#).
- See a [quick reference of Mockito](#) syntax and operations.

1b/ Instead of the JUnit core asserts, you may use the [Hamcrest library](#) to create more human-readable assertions. Consider using this library in the previous example, in particular, `assertThat()`, `is()`.

2/ Consider an application that needs to perform reverse geocoding to find a zip code for a given set of GPS coordinates. This service can be obtained in the Internet (e.g.: using the [MapQuest API](#)).

- Create the objects represented in Figure 1. `TqsHttpClient` represents a service to initiate HTTP requests to remote servers. **You don't need to implement `TqsHttpBasic`**; in fact, you should provide a substitute for it.
- Consider that we want to verify the `AddressResolver#findAddressForLocation`, which invokes a remote geocoding service, available in a REST interface, passing the site coordinates. Which is the service to fake?
- To create a test for `findAddressForLocation`, you will need to know the exact response of the geocoding service for a sample request. Assume that we will use the [MapQuest API](#). Use the browser or an HTTP client to try some samples so you know what to test for ([example 1](#)).
- Implement a test for `AddressResolver#findAddressForLocation` using a mock.
- Besides de “success” case, consider also testing for alternatives (e.g.: invalid coordinates should raise an exception).

This [getting started project](#) can be used in your implementation.

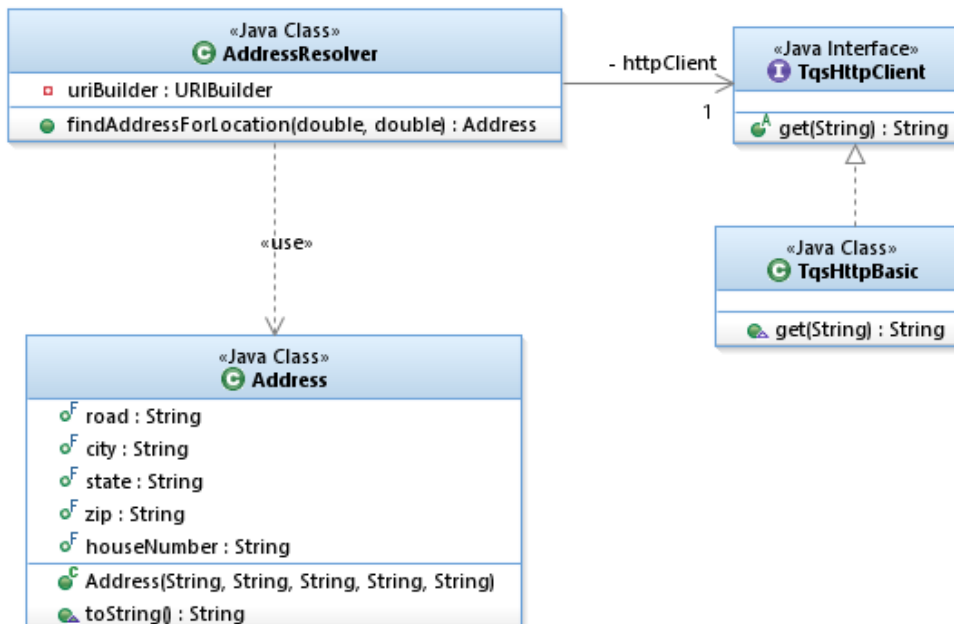


Figure 2: Classes for the geocoding use case.

3/ Consider you are implementing an integration test, and, in this case, you would use the real implementation of the module, not the mocks, in the test.

Create new test class and be sure its name end with “IT” (e.g.: `GeocodeTestIT`).

Copy the tests from the previous exercise into this new test class.

Remove all support for mocking (no dependencies on Mockito imports).
Correct the test implementation, so it used the real module.

If the “failsafe” maven plugin is configured, you should get different results with:

```
$ mvn test
```

```
$ mvn package failsafe:integration-test
```

Explore

JUnit 5 [cheat sheet](#).

Lab 3: Functional testing with web automation

Prepare

- [Selenium and separation of concerns](#)

Key Points

- Acceptance tests (or functional test) exercise the user interface of the system, as if a real user was using the application. The system is treated as a black box.
- Browser automation (control the browser interaction from a script) is an essential step to implement acceptance tests on web applications. There are several frameworks for browser automation (e.g.: Puppeteer); for Java, the most used framework is the WebDriver API, provided by Selenium (that can be used with JUnit or TestNG engines).
- The test script can easily get “messy” and hard to read. To improve the code (and its maintainability) we could apply the Page Objects patterns.
- Web browser automation is also very handy to implement “smoke tests”.

Lab

Selenium works with multiple browsers but, for sake of simplicity, the samples will be discussed with respect to Chrome/Chromium; you may adapt for Firefox.

Suggested setup:

Install Chrome/Chromium in your system (if needed), using the default installation paths.

Download the [ChromeDriver](#) and make sure it is available available in the PATH. ([GeckoDriver](#) for Firefox)

Install the “Selenium IDE” browser plugin. (or, alternatively, the Katalon Recorder).

In this lab:

1. Create a web automation with Selenium IDE recorder
2. Run the test as a Java project (JUnit 5 + Selenium)
3. Use the Web Page Object pattern

1/ Create a web automation with Selenium IDE recorder

Access the Redmine demo site and create a temporary account (<http://demo.redmine.org>)

Be sure to logout before recording the test.

Create (record) an automation macro with the Selenium IDE recorder tool to test login (a [quick start for Katalon](#) is available, which is similar to Selenium IDE):

- a) Open <http://demo.redmine.org>
- b) Sign-in with your credentials
- c) Assert that you have successfully logged in (by verifying the presence of the username)
- d) Logout

... and Stop recording. Test your macro (replay).

Add a new step, at the end, to confirm that, after logout, the home shows the “Sign in” option present. Enter this assertion “manually” (in the editor, but not recording).

2a/ Run the test as a Java project (JUnit 5, Selenium)

Prepare a (new) project to run JUnit tests and Selenium ([sample POM.xml](#) available; [alternative site](#)).

Take note of the information [in this page](#) under “quick reference”; then, in the section “Local browsers”, pick the example that suites your setup and run the test (as you usually do with JUnit).

You will have to deploy the WebDriver implementation (binary) for you browser [→ [download browser driver](#)]. Be sure to include in the system PATH.

2b/ Export and run the test (Webdriver)

Export the test from Selenium IDE into a Java test class and include it in the previous project.

Refactor the code that was generated to be compliant with JUnit 5 and the [Selenium-Jupiter extension](#).

Run the test (programmatically).

3/ Use the Web Page Object pattern

Consider the [example discussed here](#).

Implement the “Page object pattern” for a cleaner and more readable test, as suggested.

Notes:

Execute the suggested steps to interactively record the test case. Export it to Java and run the test with JUnit automation, refactoring for JUnit 5.

The text in the tutorial is somewhat old. You may **need to adapt to the current implementation** of the site under test (e.g.: IDs of page elements,...).

You should **stop** at “Increased readability” with Cucumber.

Explore

- [Puppeteer](#) - a Node library which provides a high-level API to control headless Chrome/Chromium.
- Another, more recent, [Page Object Model example](#).

Lab 4: Behavior-driven development (Cucumber in Java)

Key Points

- The [Cucumber framework](#) enables the concept of “executable specifications”: with Cucumber we use concrete examples to specify what we want the software to do. **Scenarios are written before production code.**
- Cucumber executes features (test scenarios) written with the [Gherkin language](#) (readable by [non-programmers too](#)).
- The steps included in the feature description (scenario) must be mapped into Java test code by annotating test methods with matching “expressions”. [Expressions](#) can be (traditional) regular expressions or the (new) Cucumber expressions.

Lab

Bear in mind that the integration of JUnit 5 and Cucumber is still very recent and most of the samples available in the internet are still based on JUnit 4.

In this lab you will:

- 1/ Create a simple cucumber-enabled project
- 2/ Book search example
- 3/ Integrate Cucumber with Selenium Webdriver

1/ Create a simple cucumber-enabled project

This example assumes that you have Maven available in the path and you can [invoke mvn from the command line](#). Check with:

Go through the [Cucumber getting started tutorial](#) for Java.

The example uses the command line and the IntelliJ IDEA (you may use other IDE, but you will need to adapt the instructions).

Note that:

- Start the project from the suggested Maven archetype `io.cucumber:cucumber-archetype`
- You need a test class that activates the Cucumber runner; that is the purpose of the “RunCucumberTest” file automatically included in this archetype.
- the features are stored under `src/test/resources/` and the folder structure (under [resources](#)) must mirror the package hierarchy under `src/test/java/`. In this sample: `src/test/resources/hellocucumber/` and `src/test/java/hellocumcucumber/`
- IntelliJ recognizes the .feature file type. You can even select “Run all features” to run the Cucumber tests.

2/ Book search example

To get into the “spirit” of BDD, partner with a colleague, and jointly write a couple of features to verify a book search user story. Consider a few search options (by year, etc).

Take the approach discussed in [this example](#), and write your own tests. Feel free to add different scenarios/features.

The sample uses Cucumber v2, but you should write the test steps using Cucumber 3x. Some changes are required:

In the “book search” story:

- You will need to change the parameters placeholder in the steps definition. Prefer the “cucumber expressions” (instead of regular expressions). [→ [partial snippet](#)]
- migrate the date formatter option, by creating a new datatype in a “Configurer” class, which should be placed in the same package as the test steps [→ [possible solution snippet](#)]. This defines a new custom parameter type (“date_iso_local_date_time”)
- The dates in the feature description need also to match the date mask used (aaaa-mm-dd).

In the “salary” story:

- Adapt the data table definition for the Salary use case [→ [possible solution snippet](#)]. This allows to extract a List<Employees> from the feature definition and use it as a parameter.

```
public void
the_salary_management_system_is_initialized_with_the_following_data(final
List<Employee> employees)
```

Notes:

→ you may build on the previous project or start with a new one, but stick with the base archetype (io.cucumber:cucumber-archetype).

3/ Integrate Cucumber with Selenium Webdriver

Implement the brief sample using [cucumber and selenium webdriver](#) to declare expressive web automation tests.

You may now extend this strategy to revisit the example (Weather forecast) of lab 3.

Lab 5: Multi-layer application testing with Spring Boot

Prepare

This lab is based on Spring Boot. Most of students already used the Spring Boot framework (in IES course).

If you are new to Spring Boot, then you need to develop a basic understanding or collaborate with a colleague. [Learning resources](#) are available at the Spring site.

Key Points

- Isolate the functionality to be tested by limiting the context of loaded frameworks/components. For some use cases, you can even test with just standard unit testing.
- @SpringBootTest annotation loads whole application, but it is better to limit Application Context only to a set of spring components that participate in test scenario
- @DataJpaTest only loads @Repository spring components, and will greatly improve performance by not loading @Service, @Controller, etc.
- Use @WebMvcTest to test rest APIs exposed through Controllers. Beans used by controller need to be mocked.

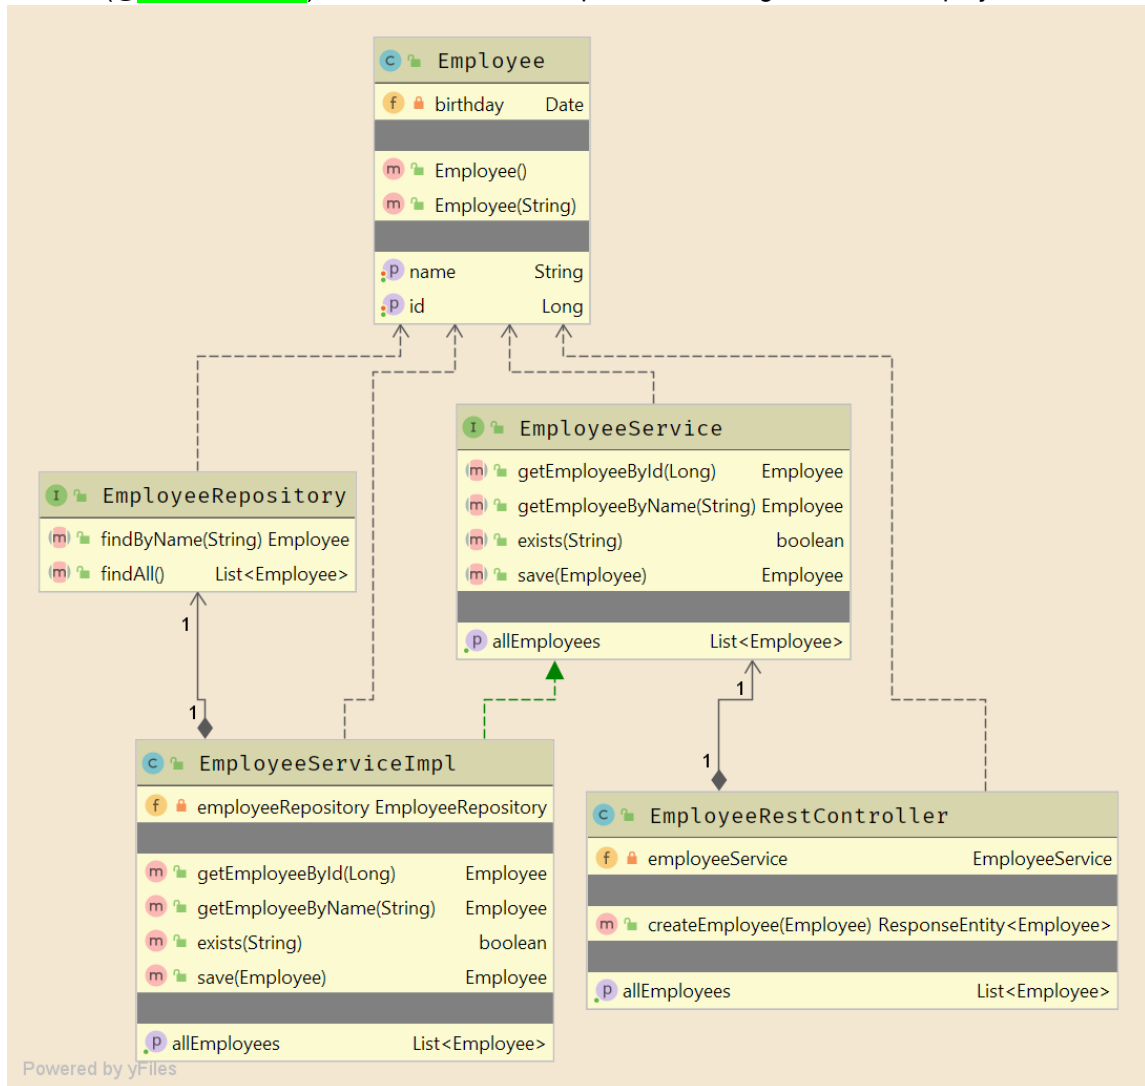
Lab

1/

Study the example available concerning a simplified [Employee management application](#) (gs-employee-manager).

This application follows commons practices to implement a Spring Boot solution:

- Employee: entity (@Entity) representing a domain concept.
- EmployeeRepository: the interface (@Repository) defining the data access methods on the target entity, based on the framework JpaRepository. “Standard” requests can be inferred and automatically supported by the framework (no additional implementation required).
- EmployeeService and EmployeeServiceImpl: define the interface and its implementation (@Service) of a service related to the “bizz logic” of the application. Elaborated decisions/algorithms, for example, would be implemented in this component.
- EmployeeRestController: the component that implements the REST-endpoint/boundary (@RestController): handles the HTTP requests and delegates to the EmployeeService.



The project contains a set of tests. Take note of the following test scenarios:

Purpose	Strategy	Notes
A/ Verify the data access services provided by the repository component. [EmployeeRepositoryTest]	Slice the test context to limit to the data instrumentation (@DataJpaTest) Inject a TestEntityManager to access the database; use directly this object to write to the database.	@DataJpaTest includes the @AutoConfigureTestDatabase. If a dependency to an embedded database is available, an in-memory database is set up. Be sure to include H2 in the POM.
B/ Verify the business logic associated with the services implementation. [EmployeeServiceImplUnitTest]	Can be achieved with a unit test, if we mock the repository behavior. Rely on Mockito to control the test and to set expectations and verifications.	Relying only in JUnit + Mockito makes the test a unit test, much faster than using a full SpringBootTest. No database involved.

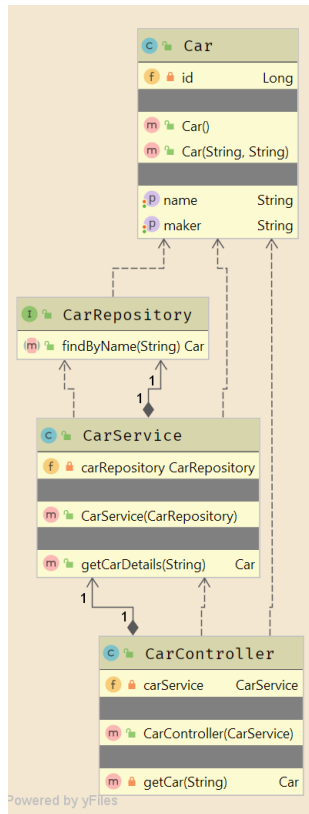
C/ Verify the boundary components (controllers). No need to test the real HTTP-REST framework; just the controller behavior. [EmployeeControllerIT]	Run the tests in a simplified light environment, simulating the behavior of an application server, by using @WebMvcTest mode. Get a reference to the server context with @MockMvc. To make the test more localized to the controller, you may mock the dependencies on the service (@MockBean); the repository component will not be involved.	MockMvc provides an entry point to server-side testing. Despite the name, is not related to Mockito. MockMvc provides an expressive API, in which methods chaining is expected.
D/ Verify the boundary components (controllers). Test the REST API on the server-side; no API client involved. [EmployeeRestControllerIT]	Start the full web context (@SpringBootTest, with Web Environment enabled). The API is deployed into the normal SpringBoot context. Use the entry point for server-side Spring MVC test support (MockMvc).	This would be a typical integration test in which several components will participate (the REST endpoint, the service implementation, the repository and the database).
E/ Verify the boundary components (controllers). Test the REST API with explicit HTTP client involved. [EmployeeRestControllerTemplateIT]	Start the full web context (@SpringBootTest, with Web Environment enabled). The API is deployed into the normal SpringBoot context. Use a REST client to create realistic requests (TestRestTemplate)	Similar to the previous case, but instead of assessing a server entry point for tests, start a API client (so request and response un/marshaling will be involved).

Review questions:

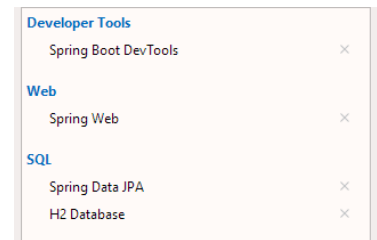
- Identify a couple of examples on the use of AssertJ expressive methods chaining.
- Identify an example in which you mock the behaviour of the repository (and avoid involvind a database).
- What is the difference between standard @Mock and @MockBean?
- What is the role of the file "application-integrationtest.properties"? In which conditions will it be used?

2/

Consider the case in which you will develop an API for a car information system.



Implement this scenario, as a Spring Boot application. Consider using the [Spring Boot Initializr](#) to create the new project (integrated in IntelliJ or online); add the “starters” for Developer Tools, Spring Web, Spring Data JPA and H2 Database.



Take the structure modeled in the classes diagram as a (minimal) reference.

In this exercise, try to force a TDD approach: write the test first; make sure the project can compile without errors; defer the actual implementation of production code as much as possible. This will be forced if we try to write the tests in a top-down approach: start from the controller, than the service, than the repository.

- Create a test to verify the CarController (and mock the CarService bean). Run the test.
- Create a test to verify the CarService (and mock the CarRepository). This can be a standard unit test with mocks.
- Create a test to verify the CarRepository persistence. Be sure to include a in-memory database dependency in the POM (e.g.: H2).
- Having all the previous tests passing, implement an integration test

to verify the API. Suggestion: use the E/ approach discussed for the previous (Employees) example.

- Adapt the integration test to use a real database. E.g.:
 - Run a mysql instance and be sure you can connect (for example, using a Docker container)
 - Change the POM to include a dependency to mysql;
 - Add the connection properties file in the resources or the “test” part of the project (see the [application-integrationtest.properties](#) in the sample project)
 - Use the `@TestPropertySource` and deactivate the `@AutoConfigureTestDatabase`.

Lab 6: Static Code analysis with Sonar Qube

Prepare

You will find a lot of demos in Youtube under for static code analysis in Java with SonarQuebe. It is not required but you may choose to have a look.

Key Points

Static code quality can be inspected to obtain quality metrics on the code base. These metrics are based on the occurrence of known weaknesses, a.k.a. “code smells”. The code is not executed (thus static analysis).

Key measures include the occurrence of problems likely to produce errors, vulnerabilities (security/reliability concerns) and code smells (bad/poor practice or coding style); coverage (ratio tested/total); and code complexity assessment.

The estimated effort to correct the vulnerabilities is called the **technical debt**. Every software quality engineer must use specialized tools to obtain realistic technical debt information.

Lab

In this lab:

Task 1: Analyze an existing project (maven-based)

Task 2: Include tests and coverage

Task 3: Define and apply quality gates

Task 1: Analyze an existing project (maven-based)

a/ [Install the SonarQube server](#) either as a local service, or by using the docker image (see instructions in the link).

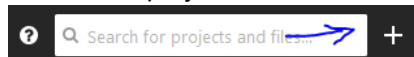
For the purpose of this lab, you don't need to configure a production database (the embedded H2 database is used by default). If, however, you want a production-like setup, you should consider using a persistent database (for example, like [this configuration](#); not required for the class).

b/ Configure the environment for [Sonar Scanner for Maven](#). Be sure to adapt the <sonar.host.url> for your setup (e.g.: <http://127.0.0.1:9000>)

c/ Select a Maven-based, Java application project to use. You may reuse one from previous labs, for example, the Euromillions from Lab 1 (part 2b), with tests passing.

Configure the POM to integrate the Sonar scanner plugin (section #2, in [this guide](#))

d/ Create a project in the Sonar Qube dashboard (default : <http://127.0.0.1:9000>).



Be sure to define the **project key equal to the maven GroupId:ArtifactId** elements.

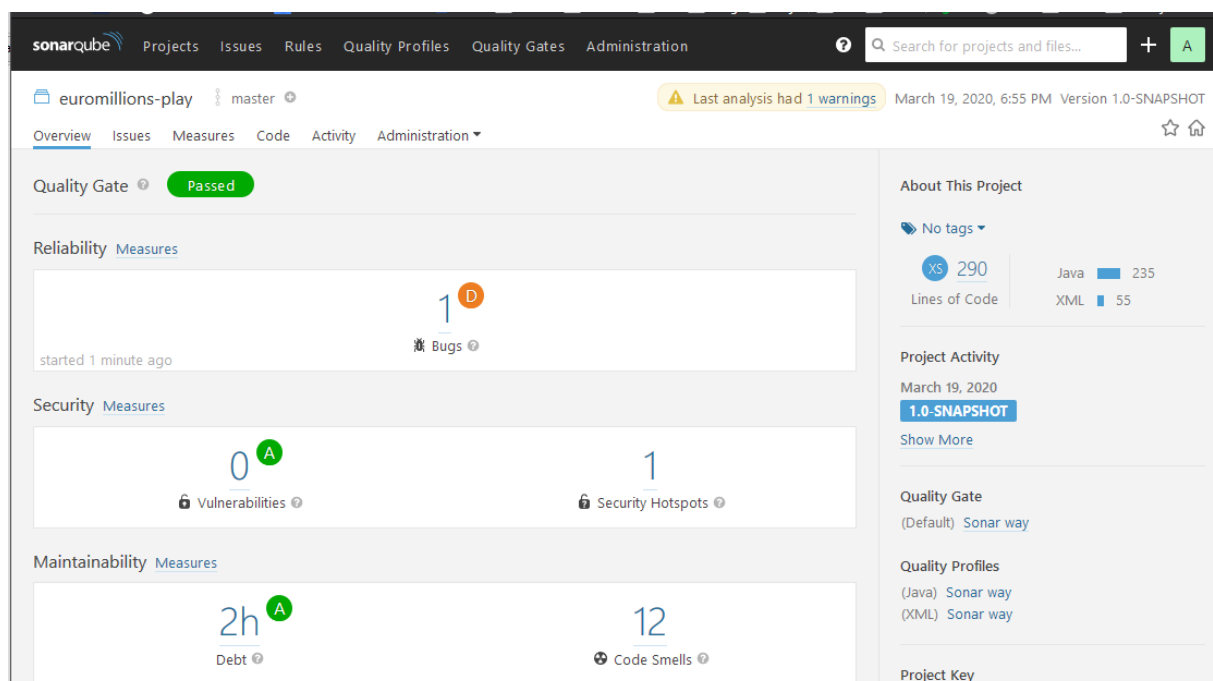
Then, generate a **token** for the project and take note for later use.

Then, in the command line, run the code analysis (adapt as needed):

```
$ mvn sonar:sonar -Dsonar.projectKey=tqslabs:euromillions-play -
```

```
Dsonar.host.url=http://127.0.0.1:9000 -Dsonar.login=9d0ca979aff1cacd052c669072275f70ba2bd512
```

e/ observe the Sonar dashboard. Has the project passed the defined quality gate?



Explore the analysis results and complete with a few sample issues (if applicable):

Issue (be severity)	Problem description	How to solve
Bug
Vulnerability (severe)		
Code smell (severe)		

f/ Describe the [metrics](#) assessed in the Quality Gate that is being used (Sonar Dashboard --> top menu --> Quality Gates view):

Metric	Conditions	Explanation (in your own words)
Coverage on New Code	Fails if < 80	How much of the UPDATED source code has been covered by the unit tests? Mixes Line coverage and Condition coverage.
...		

Task 2: Include tests and coverage

For this part, be sure you are using a project with JUnit tests implemented and passing. Let us assume the Euromillions project.

a/ Take note of the technical debt found. Explain what this value means.

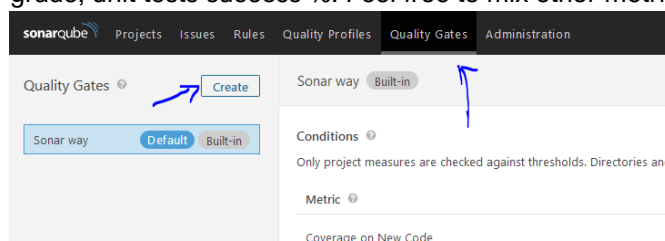
b/ Analyze the reported problems and be sure to **correct the severe** code smells reported (critical and major).

c/ Code coverage reports require an additional plugin. Be sure to use a project with unit tests and configured code coverage (e.g.: [add the jacoco plugin](#) to maven). You may have already did it in Lab 2 (2d).

d/ Run the static analysis and observe/explore the coverage values on the Sonar Qube dashboard. How many lines are “uncovered”? And how many conditions?

Task 3: Define and apply quality gates

a/ Define a custom [quality gate](#) to this project. Consider using: coverage %, Maintainability Rating grade, unit tests success %. Feel free to mix other metrics.



b/ Add an increment to the source code. You may try to introduce some “bad smell”.
Does the updated project pass the quality gate?

Explore:

- Use the SonarQube to inspect a project of your own, maybe from another course. Note that free analyzers are not available for all languages.
- A related [tutorial by Baeldung](#).
- public projects on [Sonar cloud](#) that you can browse and learn.

Lab 7: Continuous integration with Jenkins

Prepare

You should already have Java JDK, Maven and Docker installed and configured in your environment. If not available, be sure to [Install Docker CE](#) for your platform. For Linux, be sure to configure Docker to [run as non-root user](#).

You may run Jenkins in one these (main) scenarios:

- A. Jenkins is installed as a regular application/service in your machine/server. The Jenkins workspace is in the regular filesystem and Jenkins can access the machine's tools and environment. Jenkins may invoke Docker components to complete some stages.
- B. Jenkins is running in a container, and the agents used to run the jobs are themselves Docker containers, with proper configuration. Sharing volumes between containers will avoid redundant transfers.

In this lab, we will run Jenkins in scenario A and then in scenario B.

Lab

Tasks 1: Use Jenkins as standalone application and create jobs by configuration

This would be the “old way” of configure Jenkins, but it prides a good understanding of the tool.

You will also need:

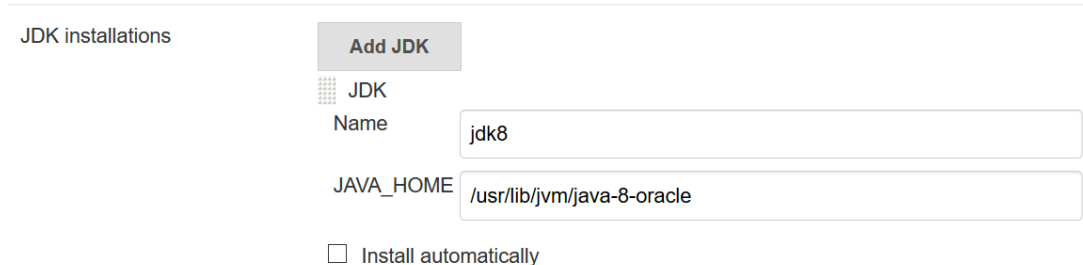
- A Git repo with a Maven Java project, ready to run (can be a local repo, in the file system).
- a) Start by getting Jenkins. Use the [jenkins.war LTS release](#) (generic java package, last option in the download page). Run Jenkins in the terminal:


```
$ java -jar jenkins.war --httpPort=8080
```

Note1: take note of the auto-generated key printed in the terminal. You will need it to activate Jenkins.

Note2: port 8080 should be available, otherwise change the port.
 - b) Verify that Jenkins server is running: <http://localhost:8080>
Complete the post install steps including:
 - unlocking Jenkins
 - install the suggested plugins (may take a while...).
 - Create the admin user

- c) Configure Java and Maven
- Access Jenkins > Manage Jenkins > Global Tool Configuration
- Add/Configure JDK.
- Uncheck the automatic install and point to a local JDK.
- Add/Configure Maven
- Uncheck the automatic install and point to a local maven.

JDK


JDK installations

Add JDK

JDK

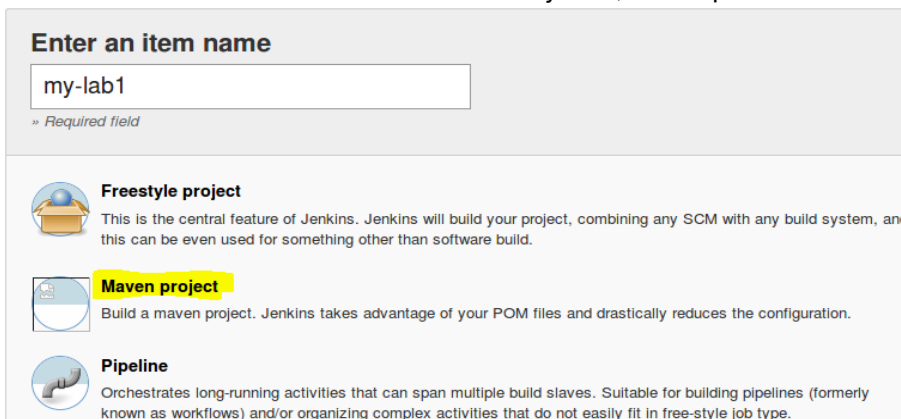
Name

JAVA_HOME

☐ Install automatically

- d) Install the “maven integration” and “pipeline maven integration” plugin (Manage Jenkins > Plugins > Available).
- e) Let's create a **new job**

Create a new project (New item). This should be a “Maven Project” item. Note that the project name will be used to create a folder in the filesystem; avoid special characters and spaces.



Enter an item name

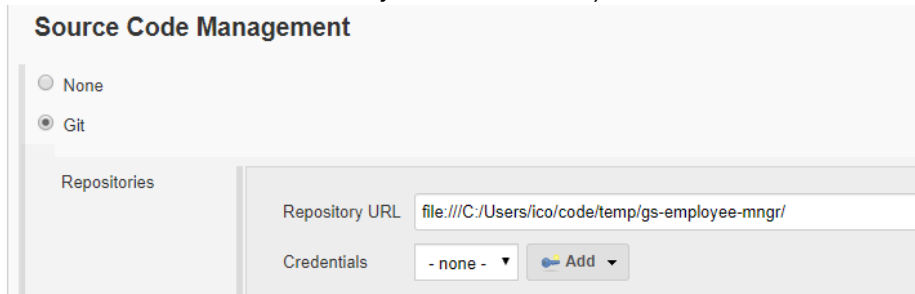
» Required field

Freestyle project
This is the central feature of Jenkins. Jenkins will build your project, combining any SCM with any build system, and this can be even used for something other than software build.

Maven project
Build a maven project. Jenkins takes advantage of your POM files and drastically reduces the configuration.

Pipeline
Orchestrates long-running activities that can span multiple build slaves. Suitable for building pipelines (formerly known as workflows) and/or organizing complex activities that do not easily fit in free-style job type.

In the Source code management, select Git. Configure the local repository URL (or a remote and add the credentials to the Jenkins keychain, as needed).



Source Code Management

☐ None

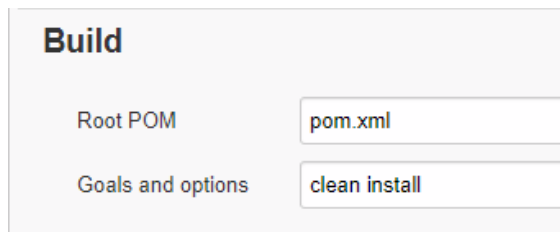
☒ Git

Repositories

Repository URL

Credentials **Add**

Configure the build step to adequate maven goals. Provide the path to the local pom.xml in “Root POM” field.

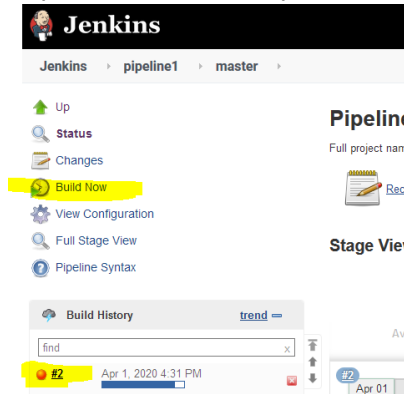


Build

Root POM:

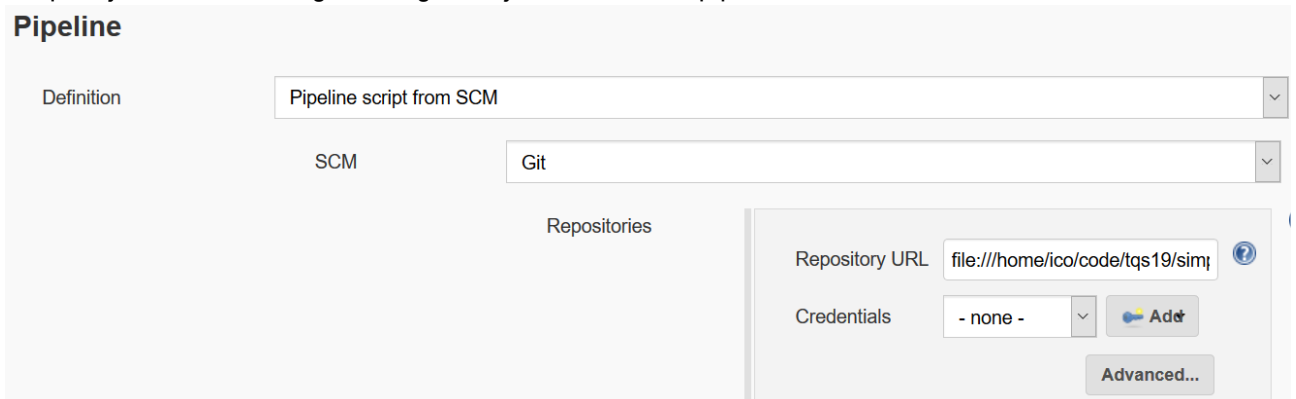
Goals and options:

- f) Manually request a build. Verify the build results.
- g) Inspect the console output



Tasks 2: Create a simple pipeline for a Maven project and poll the SCM (no Dockerization)

- h) Configure a new job, for a **pipeline project** hosted on a git repository. (you may reuse the same repository)
- i) Specify that the build triggers are given by a “Jenkinsfile” pipeline definition file.



The screenshot shows the Jenkins Pipeline configuration screen. The 'Definition' dropdown is set to 'Pipeline script from SCM'. The 'SCM' dropdown is set to 'Git'. The 'Repositories' section shows a list of repositories. The 'Repository URL' field is set to 'file:///home/ico/code/tqs19/sim'. The 'Credentials' dropdown is set to '- none -'. There is an 'Add' button next to the credentials dropdown and an 'Advanced...' button at the bottom right.

- j) Create a “Jenkinsfile” in the root of the project, with the following content:

```

pipeline {
    agent any
    tools {
        jdk 'jdk8'
        maven 'mvn3'
    }
    stages {
        stage('test java installation') {
            steps {
                sh 'java -version'
            }
        }
        stage('test maven installation') {
            steps {
                sh 'mvn -version'
            }
        }
    }
}

```

→ Note: the name of the tools (“jdk8”, “mvn3”) must match the same names you defined previously in the global tool configuration.

→ Note: if the build is to apply in a subfolder, you need to switch to that folder in the beginning of the step. E.g.:

```
steps {
    dir('your-sub-directory'){
        sh "$PWD"
    }
}
```

Try to run the jenkins job (**after committing** the Jenkins file).

Inspect the console output of the Job. You should confirm that the Job ran and printed the version information of the tools (nothing fancy...).

- k) Configure an automated polling strategy (Build triggers -> SCM Polling) instead of manual. See the help button near the text box (in the Jenkins interface) for syntax and examples.

- l) Update the Jenkins file to include actual build of the project:

```
pipeline {
    agent any
    tools {
        jdk 'jdk8'
        maven 'mvn3'
    }
    stages {
        stage('Install') {
            steps {
                sh "mvn clean install"
            }
            post {
                always {
                    junit '**/target/*-reports/TEST-*.xml'
                }
            }
        }
    }
}
```

Commit the changes. Observe in detail the console output on the building process.

- m) Explore the project workspace (the folder used by jenkins) -> [Workspaces](#)
- n) Add more changes locally and push them to the mainline:
- introduce a failing test (e.g. with [fail\(\)](#))
 - commit to the main line
 - confirm that SCM is pooled and the build fails (changes from “stable” to “with errors”)

Task 3: Create a Maven-based pipeline and use Docker containers

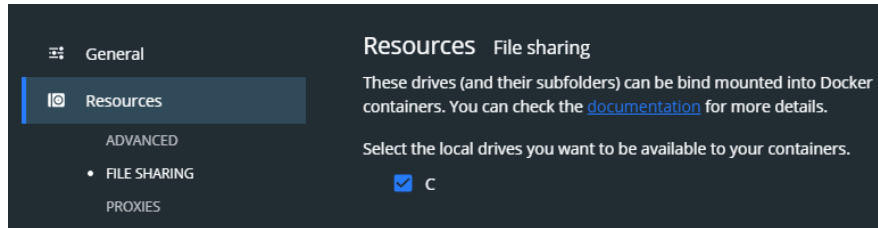
In this task, you will need to use Docker.

Take the Jenkins' tutorial on [Building a Java Maven app](#).

In this example, two containers are used for (1) running the Jenkins service and (2) implement the build stages.

Notes:

- You may need to activate the option “Share drive” if you are running Docker in Windows. (specially in the step nr. 4 of the tutorial)



- To open a shell inside a container:
`docker container exec -it jenkins-blueocean bash`

For example, to get the hash to unlock the Jenkins installation:
`bash-4.4$ cat /var/jenkins_home/secrets/initialAdminPassword`

- If you get errors checking-out from a local repository, be sure to use a remote, online Git repository instead.

Lab 8: Code review in the context of pull/merge requests

Prepare

You will work in informal teams. You should find at least another colleague to pair-with; better if you “reuse” a larger team and an existing project (e.g.: the same team and project from “Projeto em Informática”).

Lab

In this exercise, you will **demonstrate the use of code reviews by peers**.

You should collect evidence of the work done (screenshots, tool dashboards,...) to submit a brief report.

Task 1:

Apply the [GitHub flow](#) to a project of your own (or other that is based on [feature-branching](#)). You may use a project under progress in another course (if you are already using feature-branch in another project, use that context in this task).

Though you are encouraged to use an existing code base, you still may set up the scenario from scratch. (Does not need to be in Java.)

You will need a git repository with working code, shared in a team of (2+) contributors, in popular git hosting platform (e.g.: [GitHub](#), [GitLab](#), [Bitbucket](#)).

- Develop at least a new increment (feature) per each member in the team.
- Integrate the feature by issuing a pull request (= merge request). Be sure to **have code review interactions in the context of the pull request**. Code reviews should be blocking, meaning, another person needs to explicitly approve your pull request.

In your (brief) report:

- Describe the kind of workflow and the rules adopted by the team (Gitflow workflow? Github flow? [Rules for the commit message](#)?...).

- Show evidence of the pull requests/code review that were performed. It should be possible to the reader to appreciate the “content”/usefulness of the review (e.g.: what was comment?)
- Optional: provide a link to the repo.

The members prepare a single team report, but everyone submit it in Moodle.

Explore

Guidelines and cheat-sheets:

- Code review guidelines: [GitLab](#) guidelines (for contributors to the GitLab project)
- Coding style & rules: [Android OS project](#)
- Git cheat-sheets: [Hyperpoliglot](#) (comprehensive)| [interactive visualization](#) | git commands [full reference](#)

Lab 9: Performance testing

Key points

There are many [frameworks for performance testing](#), either hosted or on own premises. These kind of frameworks ensure: load generation; scalable architectures; metrics collection for evidence.

Lab

In this lab, we will use [JMeter framework](#) to create synthetic load and collect performance metrics.

JMeter is a Java application; just [download and unzip](#). If needed, here are more [details on JMeter installation](#).

Task 1/ Simulating Web workflow

HTTP workflows usually include configuration elements, such as clearing cookies and cache; issue different kinds of HTTP Requests; create explicit delays to simulate usage flows.

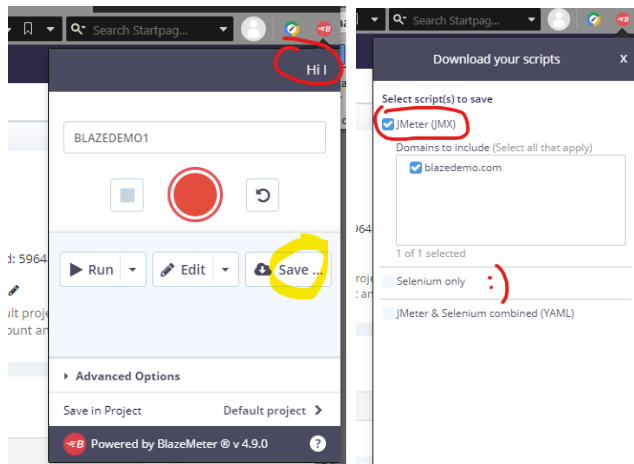
Follow the examples in this [guided exercise](#).

- be moderate on the number of users (ThreadGroup)
- add a Response Time Graph to monitor the performance trends.
- feel free to adapt/experiment with your own projects.

Note 1: as illustrated in the “Your Third JMeter Script” section of the tutorial, you may use the [BlazeMeter plugin for Chrome](#) to interactively record a test and export it to the JMeter format. This is very convenient to capture a workflow.

You should have an account and be logged-in (with BlazeMeter) to have access to the option to save the JMX test.

Open the recorded teste in JMeter IDE and inspect it. Note the use of (constant) Timers to have the same delays that happened when you recorded the test.



Note 2: the JMeter IDE should be used for the test development and debugging. For production, you should run the tests from the command line (or, better, with a cloud-based tool).

Task 2 Testing a REST endpoint

Follow the example available in this [guided tutorial](#). You will need an account for the OctoPerf service (once a student has one, you may share it with a few colleagues).

- Skip step 3.3 (uses a deprecated plugin; does not influence the result.)
- Step 4.1: adapt the JsonPath expression to work with the results that you are obtaining from the API. E.g.: `$.[0]['name']` is equal to "Default".
- Skip step 5 (it would be interesting if you add several Workspaces in the platform.)

Feel free to use the REST API testing in your own projects (e.g.: for the mid-term assignment).

Task 3 Database load [optional]

See also: related sample, with [tutorial steps](#)

Databases are a critical element in enterprise applications and often requiring optimization strategies. Implement the following test on a database of yours (or setup a new instance):

- a) Copy the appropriate JDBC driver to the JMeter lib folder
 - b) Create a new Test Plan and add users (ThreadPool).
 - note: you may speed up by using the JDBC template (Menu File → Templates → JDBC)
 - c) Add the connection pool configuration (JDBC Connection Configuration)
 - d) Add/edit the JDBC request properties:
 - a. Bound to the previous connection pool (the same name used in the JDBC Connection Configuration)
 - b. Enter a SQL Query
 - e) Add a Timer based on a Poisson distribution
 - f) Add a listener and view results in a table and a graph
- Run and verify the output, with different workloads.

Explore

JMeter can be invoked [from Maven](#).

[JMeter tests can be integrated in the Jenkins](#) CI/CD pipeline.